
MLPR Exam Project: Gender Detection

Mattia Rosso [s294711]

September 4, 2022

This project is intended to show a binary classification task on a dataset made of 12 continuous observations coming from speaking embeddings. A speaker embedding represents a small-dimensional, fixed size representation of an utterance. Features can be seen as points in the m-dimensional embedding space (and the embeddings have already been computed). This is a task where classes are balanced both in training and evaluation set.

1 Dataset analysis

1.1 Training and evaluation sets

The datasets provided are:

- Training Set: 3000 samples belonging to Male class (Label = 0) and 3000 samples belonging to Female class (Label = 1).
- Evaluation Set: 2000 samples belonging to Male class (Label = 0) and 2000 samples belonging to Female class (Label = 1).

We will use the Training Set to perform all the analysis and only once we will have selected the most promising models we will train these ones using the entire Training Set and final considerations will be done considering their behaviour on the Evaluation Set.

1.2 Features Statistics

All the features are contiguous and their main statistics can be showed through a boxplot in figure 1.

1.3 Z-normalization

A useful operation that can be applied in order to avoid to deal with numerical issues and to make data more

uniform is to apply Z-normalization as a preprocessing step. We are going to transform each sample of our dataset as:

$$z_{i,j} = \frac{x_{i,j} - \mu_j}{\sigma_j} \forall x_i \in D$$

Where $z_{i,j}$ is the Z-normalized value corresponding to the feature j of sample i while μ_j and σ_j are, respectively, the mean and the variance computed over all the values for feature j .




Figure 1: Raw features boxplots

1.4 Features distribution

By plotting one histogram for each feature, separated for classes male and female, it is possible to show if the samples (separately for each feature) follow a Gaussian distribution and how well. This is done in order to understand whether a pre processing step like Gaussianization can be useful or not for our training data.

We can notice from figure 2 that almost all the features are already well-distributed (w.r.t. the Gaussian distribution) except for features 3, 7, 9. We can thus apply an additional pre-processing step in order to

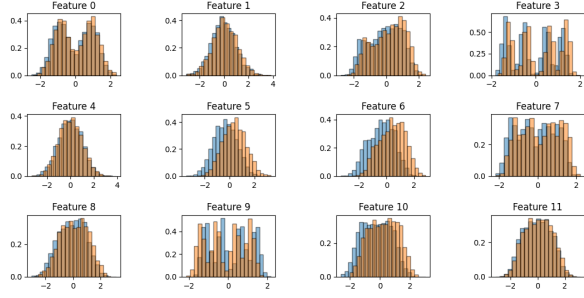


Figure 2: Z-normalized features distribution

make all the features following a Gaussian distribution: this step is called Gaussianization and it will be always applied after Z-Normalization preprocessing.

1.4.1 Gaussianization

Gaussianization is a pre-processing step that maps each feature to values whose empirical cumulative distribution are well approximated by a Gaussian cumulative distribution function. For each feature x that we want to gaussianize we firstly compute the rank over the dataset:

$$r(x) = \frac{\sum_{i=1}^N \mathbb{I}[x_i \leq x] + 1}{N+2}$$

where \mathbb{I} is the indicator function (1 when the condition inside $[\]$ is true, 0 otherwise). Actually, we are counting how many samples in the dataset D have a greater value with respect to the feature we are computing the rank on.

The next step is to compute the transformed feature as $y = \Phi^{-1}(r(x))$ where Φ is the inverse of the cumulative distribution function.




Figure 3: Gaussianized features distribution

1.5 Features correlation

We can show how much features are correlated by using a heatmap plot showing a darker color inside

cells $[i, j]$ for which it exists an high correlation among feature i and feature j . We are going to use the Pearson correlation coefficient to compute the correlation among feature X and feature Y :

$$\frac{Cov(X,Y)}{\sqrt{Var(X)}\sqrt{Var(Y)}}$$




Figure 4: Z-normalized features correlation: grey the whole dataset, orange F class, blue M class

Correlation is quite high among most of the features. We can understand that applying PCA would be meaningful for values of m not below 11 or 10 at most. For smaller values of m we would probably loose important information coming from high-correlated features.

2 Dimensionality reduction

Before proceeding with the classification task we will spend some words on the possible dimensionality reduction techniques that we have analyzed and that could be applied: PCA and LDA.

2.1 PCA

As already anticipated, given the heatmap in figure 4 we can observe that PCA with reasonable values of m can be applied. PCA is a dimensionality reduction technique that, given a centered dataset $X = \{x_1, \dots, x_k\}$, it aims to find the subspace of \mathbb{R}^n that allows to preserve most of the information (the directions with the highest variance).

Starting from the sample covariance matrix

$$C = \frac{1}{K} \sum_i (x_i - \bar{x})(x_i - \bar{x})^T$$

we compute the eigen-decomposition of $C = U\Sigma U^T$ and project the data in the subspace spanned by the m columns of U corresponding to the m highest eigenvalues:

$$y_i = P^T(x_i - \bar{x})$$

where P is the matrix corresponding to the m columns of U associated to the m highest eigenvalues of C . In this preliminary analysis, in order to select the optimal m , we can use a cross-validation approach by inspecting how much of the total variance of the data we are able to retain by using different values for m .

We exploit the fact that each eigenvalue corresponds to the variance along the corresponding axis and the eigenvalues are the elements of the diagonal of the matrix Σ . We selecte m as:

$$\min_m s.t \frac{\sum_{i=1}^m \sigma_i}{\sum_{i=1}^n \sigma_i} \geq t$$




Figure 5: 3-fold cross validation for PCA impact evaluation

We can clearly understand from figure 5 that values of $m < 9$ would rapidly decrease the amount of retained variance of the data. We will see better later on how choosing a too small value of m would badly impact the performances of the different classifiers.

2.2 LDA

PCA technique is unsupervised so we have no guarantee of obtaining discriminant directions. Despite the fact that LDA allows to find at most $C - 1$ discriminant directions (where C is the number of classes), so it makes no sense to apply it as a dimensionality reduction technique in a binary classification task, it can be used as a linear classifier and we can understand it by its definition; LDA maximizes the between-class variability over the within-class variability ratio for the transformed samples:

$$\mathcal{L}(w) = \frac{s_B}{s_W} = \max_w \frac{w^T S_B w}{w^T S_W w}$$

It can be proved that the optimal solution corresponds to the eigenvector of $S_W^{-1} S_B$ corresponding to the largest eigenvalue. Once that we have estimated w we can project the test samples over w and assign the class label looking at the score obtained:

$$C(x_t) = \begin{cases} C_1, & \text{if } w^T x_t \geq t. \\ C_0, & \text{if } w^T x_t < t. \end{cases}$$

It will be proved later how this model is related to the Tied Gaussian Generative Classifier model.

3 Classification models analysis

3.1 Premises

In the next paragraphs we are going to compare different classification models. We will employ a k-fold cross validation technique (with $k = 3$) for model evaluation. We will consider three types of applications:

$$\begin{aligned} (\tilde{\pi}, C_{fp}, C_{fn}) &= (0.1, 1, 1) \\ (\tilde{\pi}, C_{fp}, C_{fn}) &= (0.5, 1, 1) \\ (\tilde{\pi}, C_{fp}, C_{fn}) &= (0.9, 1, 1) \end{aligned}$$

and the target application (the one we will optimize for) will be the balanced one:

$$(\tilde{\pi}, C_{fp}, C_{fn}) = (0.5, 1, 1)$$

We are interested in selecting the most promising approach and we will infact perform measures in term of minimum detection cost:

$$DCF = \frac{DCF_u(\pi_T, C_{fn}, C_{fp})}{\min(\pi_T, C_{fn}, (1-\pi_T)C_{fp})} = \frac{\pi_T C_{fn} P_{fn} + (1-\pi_T) C_{fp} P_{fp}}{\min(\pi_T, C_{fn}, (1-\pi_T)C_{fp})}$$

and for $\min DCF$ computation we will look for the threshold:

$$t' = -\log\left(\frac{\tilde{\pi}}{1-\tilde{\pi}}\right)$$

that allows us to obtain the lowest possible DCF (as if we knew in advance this optimal value for threshold).

3.2 Gaussian models

The first class of models we are going to analyze are the generative Gaussian models. We assume that, given the dataset X , the sample x_t is a realization of the R.V. X_t . A simple model consists in assuming that our data, given the class, can be described by a Gaussian distribution:

$$(X_t | C_t = c) \sim (X | C = c) \sim \mathcal{N}(x_t | \mu_c, \Sigma_c)$$

Since we are dealing with a binary classification task we will assign a probabilistic score to each sample in terms of the class-posterior log-likelihood ratio:

$$llr(x_t) = \log r(x_t) = \log \frac{P(C=h_1|x_t)}{P(C=h_0|x_t)}$$

We can expand this expression by writing:

$$\log r(x_t) = \log \frac{f_{X|C}(x_t|h_1)}{f_{X|C}(x_t|h_0)} + \log \frac{\pi}{1-\pi}$$

$$\text{with } f_{X|C}(x_t|c) = \mathcal{N}(x_t | \mu_c, \Sigma_c)$$

While the training phase consists in estimating the model parameters of the Multivariate Gaussian distribution the scoring phase consists in computing the log-likelihood ratio (first term of the equation) for each sample. It will be then compared with a threshold specific for each application to compute the $\min DCF$. What it differentiates the different Gaussian models is the way how we estimate the model parameters of the Gaussian distribution.

3.2.1 MVG Gaussian Classifier

The ML solution to the previous described problem is given by the empirical mean and covariance matrix for each class:

$$\mu_c^* = \frac{1}{N_c} \sum_{i=1}^N x_{c,i}$$

$$\Sigma_c^* = \frac{1}{N_c} \sum_{i=1}^N (x_{c,i} - \mu_c^*)(x_{c,i} - \mu_c^*)^T$$

3.2.2 Naive Bayes Classifier

The Naive Bayes assumption simplifies the MVG full covariance model stating that if we knew that for each class the components are approximately independent we can assume that the distribution $X|C$ can be factorized over its components. The ML solution to this problem is:

$$\mu_{c,[j]}^* = \frac{1}{N_c} \sum_{i|c_i=c} x_{i,[j]}$$

$$\sigma_{c,[j]}^2 = \frac{1}{N_c} \sum_{i|c_i=c} (x_{i,[j]} - \mu_{c,[j]}^*)^2$$

The density of a sample x can be expressed as $\mathcal{N}(x|\mu_c, \Sigma_c)$ where μ_c is an array where each element $\mu_{c,[j]}$ is the the mean for each class for each component while Σ_c is a diagonal covariance matrix. The Naive Bayes classifier corresponds to the MVG full covariance classifier with a diagonal covariance matrix.

3.2.3 Tied Gaussian Classifier

This model assumes that the covariance matrices of the different classes are tied (we consider only one covariance matrix common to all classes). We are assuming that:

$$f_{X|C}(x|c) = \mathcal{N}(x|\mu_c, \Sigma)$$

so each class has its own mean but the covariance matrix is the same for all the classes. The ML solution to this problem is:

$$\mu_c^* = \frac{1}{N_c} \sum_{i=1}^N x_{c,i}$$

$$\Sigma^* = \frac{1}{N} \sum_c \sum_{i|c_i=c} (x_i - \mu_c^*)(x_i - \mu_c^*)^T$$

This model is strongly related to LDA (used as a linear classification model). By considering the binary log-likelihood ratio of the tied model we obtain a linear decision function:

$$l(r(x)) = \log \frac{f_{X|C}(x|h_1)}{f_{X|C}(x|h_0)} = x^T b + c$$

where b and c are functions of class means and (tied) covariance matrix. On the other hand, projecting over the LDA subspace is, up to a scaling factor k , given by:

$$w^T x = k \cdot x^T \Lambda (\mu_1 - \mu_0)$$

where $\Lambda(\mu_1 - \mu_0) = b$. The LDA assumption that all the classes have the same within class covariance matrix is related to the assumption done for the tied model.

3.2.4 Gaussian Models Comparison

Table 1: Gaussian Models

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Z-normalized features - no PCA			
Full Cov	0.128	0.048	0.125
Tied Cov	0.122	0.046	0.127
Naive Bayes	0.822	0.567	0.856
Z-normalized features - PCA(m=11)			
Full Cov	0.265	0.100	0.231
Tied Cov	0.257	0.098	0.227
Naive Bayes	0.278	0.108	0.245
Z-normalized features - PCA(m=10)			
Full Cov	0.303	0.115	0.267
Tied Cov	0.293	0.112	0.264
Naive Bayes	0.306	0.121	0.283
Gaussianized features - no PCA			
Full Cov	0.218	0.078	0.191
Tied Cov	0.208	0.078	0.189
Naive Bayes	0.813	0.586	0.847
Gaussianized features - PCA(m=11)			
Full Cov	0.227	0.087	0.218
Tied Cov	0.215	0.084	0.208
Naive Bayes	0.278	0.106	0.257
Gaussianized features - PCA(m=10)			
Full Cov	0.223	0.084	0.211
Tied Cov	0.212	0.082	0.207
Naive Bayes	0.279	0.103	0.254

A graphical version of the table can be helpful in analyzing results:

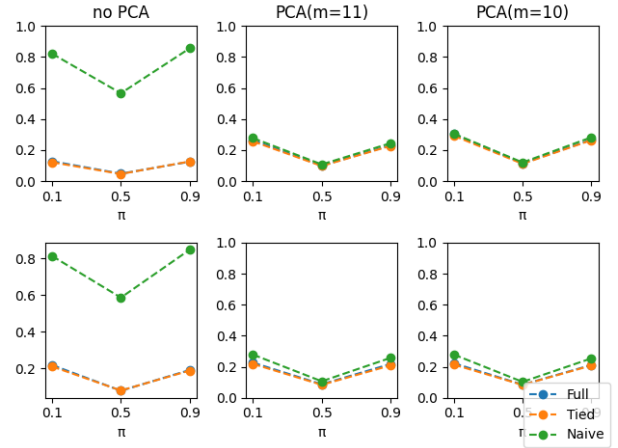


Figure 6: Top: Z-Normalized feautres, Bottom: Gaussianized features

- We can notice that Full-Cov model and Tied-Cov model achieve very good and similar results with slightly better performances for the Tied-Cov model (this means that the covariance matrices of the different classes are similar and infact the model provides more robust estimates in this case).
- Gaussianization pre processing doesn't really help

in achieving better results because data are already well distributed according to the Gaussian assumptions (this was also noticed previously while discussing the Gaussianization pre-processing).

- Naive Bayes assumption doesn't hold really well, in particular if PCA is not applied. When PCA is applied it has a really good impact only on Naive Bayes model and especially if also combined with Gaussianization pre-processing (the Naive Bayes assumption that the covariance matrix of each class is diagonal holds better with lower features dimensionality and improved features distribution from a Gaussian point of view).
- Regarding Full and Tied models with PCA(m=11) there is no high performance degradation since the minDCF values obtained are still good. For lower values of m the models become less able in taking decisions and the minDCF increases
- For the MVG classifiers best performances are achieved by the Tied-Cov classifier with only Z-normalization pre processing and without PCA. Really good performances are also achieved by the Tied-Cov model trained with Gaussianized features and no-PCA (the PCA(m=11) version of this model achieves worse but comparable result w.r.t. to the no-PCA version and it can be selected to try to better avoid overfitting by reducing the features space and also for reducing computational effort)

Best Gaussian Models:

- Tied Cov (Z-Normalization, no PCA)
- Tied Cov (Gaussianization, PCA(m=11))

3.3 Logistic Regression Classifier

Logistic Regression is a discriminative classification model. Starting from the results obtained from the Tied Gaussian classifier we consider the linear decision function obtained from the expression of the posterior log-likelihood ratio:

$$l(x) = \log \frac{P(C=h_1|x)}{P(C=h_0|x)} = \log \frac{f_{X|C}(x|h_1)}{f_{X|C}(x|h_0)} + \log \frac{\pi}{1-\pi} = w^T x + b$$

where b takes into account all the prior information. Given w and b we can compute the expression for the posterior class probability:

$$P(C = h_1|x, w, b) = \frac{e^{(w^T x + b)}}{1 + e^{(w^T x + b)}} = \sigma(w^T x + b)$$

where $\sigma(x) = \frac{1}{1+e^{-x}}$ is the sigmoid function. LR assumes that the decision rules will be hyperplanes orthogonal to w .

3.3.1 Linear Logistic Regression (LLR)

We are going to look for the minimizer of the function:

$$J(w, b) = \frac{\lambda}{2} \|w\|^2 + \frac{1}{n} \sum_{i=1}^n \log(1 + e^{-z_i(w^T x_i + b)})$$

where λ is an hyperparameter that represents the regularization term (needed to make the problem solvable in case of linearly separable classes).




Figure 7: minDCF for different values of λ and different priors

Table 2: Linear Logistic Regression - 3-fold cross validation

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Z-normalized features - no PCA			
LLR ($\lambda = 10^{-3}$)	0.170	0.059	0.155
LLR ($\lambda = 10^{-5}$)	0.132	0.047	0.127
LLR ($\lambda = 10^{-6}$)	0.132	0.047	0.126
Z-normalized features - PCA(m=11)			
LLR ($\lambda = 10^{-3}$)	0.278	0.104	0.232
LLR ($\lambda = 10^{-5}$)	0.269	0.098	0.221
LLR ($\lambda = 10^{-6}$)	0.268	0.098	0.222
Z-normalized features - PCA(m=10)			
LLR ($\lambda = 10^{-3}$)	0.299	0.113	0.263
LLR ($\lambda = 10^{-5}$)	0.297	0.114	0.263
LLR ($\lambda = 10^{-6}$)	0.297	0.114	0.262

Table 3: Best models analyzed up to now

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Gaussian Models			
Tied Cov (Z-Norm, no PCA)	0.122	0.046	0.127
Tied Cov (Gau, PCA(m=10))	0.212	0.082	0.207

- The choice of λ appear to be critical for all the applications, in particular for the unbalanced ones. By observing figure 7 it is clear that for values of λ greater than 10^{-3} the minDCF rapidly increases for all the considered applications.
- PCA never helps in achieving better results.

Comparison: the LLR model trained with $\lambda = 10^{-6}$ and no-PCA achieves really similar results with the ones achieved by the Tied-Cov Gaussian model (Z-Norm, no PCA) for $\tilde{\pi} = 0.5$ and $\tilde{\pi} = 0.9$ applications but a worse performance is obtained for $\tilde{\pi} = 0.1$. The LLR model behaves better than the Tied-Cov Gaussian model (Gau, PCA(m=10)) in all the three applications.

Selected LLR Model:

- Z-normalized features, $\lambda = 10^{-6}$, no PCA

3.3.2 Quadratic Logistic Regression (QLR)

Now we are going to train a Quadratic LR model by performing features expansion. For binary linear LR the separation surfaces are linear decision functions as already discussed (and we obtain the same form as for the Tied Gaussian classifier). By looking instead at the separation surface obtained through the MVG Gaussian classifier we have:

$$\log \frac{P(C=h_1|x)}{P(C=h_0|x)} = x^T A x + b^T x + c = s(x, A, b, c)$$

This expression is quadratic in x but it's linear in A and b . We could rewrite it to obtain a decision function that is linear for the expanded features space but quadratic in the original features space. Features expansion is defined as:

$$\Phi(x) = \begin{bmatrix} \text{vec}(xx^T) \\ x \end{bmatrix}, w = \begin{bmatrix} \text{vec}(A) \\ b \end{bmatrix}$$

where $\text{vec}(X)$ is the operator that stacks the columns of X . In this way the posterior log-likelihood is expressed as:

$$s(x, w, c) = s^T \phi(x) + c$$

We are now going to train the Linear Logistic Regression model using features vectors $\phi(x)$.




Figure 8: minDCF for different values of λ and different priors

Table 4: Quadratic Logistic Regression - 3-fold cross validation

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Z-normalized features - no PCA			
QLR ($\lambda = 10^{-3}$)	0.187	0.063	0.149
QLR ($\lambda = 10^{-5}$)	0.153	0.052	0.142
QLR ($\lambda = 10^{-6}$)	0.150	0.053	0.141
Z-normalized features - PCA(m=11)			
QLR ($\lambda = 10^{-3}$)	0.267	0.098	0.222
QLR ($\lambda = 10^{-5}$)	0.263	0.098	0.230
QLR ($\lambda = 10^{-6}$)	0.305	0.096	0.230
Z-normalized features - PCA(m=10)			
QLR ($\lambda = 10^{-3}$)	0.299	0.111	0.241
QLR ($\lambda = 10^{-5}$)	0.307	0.109	0.249
QLR ($\lambda = 10^{-6}$)	0.305	0.109	0.248

- By rapidly looking at the results obtained we can say that quadratic version of the logistic regression performs worse w.r.t the linear version (the one without features expansion).

Table 5: Best models analyzed up to now

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Gaussian Models			
Tied Cov (Z-Norm, no PCA)	0.122	0.046	0.127
Tied Cov (Gau, PCA(m=10))	0.212	0.082	0.207
Logistic Regression Models			
LLR (Z-Norm, $\lambda = 10^{-6}$, no PCA)	0.132	0.047	0.126

- The choice of λ is still critical and $\lambda \leq 10^{-5}$ still remains the best choice
- Regarding the target application we are able to reach similar results comparing to the linear version of the LR while the unbalanced applications are more penalized.
- When PCA is applied no effective improvements are obtained.

Comparison: With respect to Gaussian models comparable performances are achieved for $\lambda = 10^{-6}$ and no PCA even if the model performs slightly worse. The quadratic model performs also worse than the linear model and we won't consider it in score calibration. Selected QLR Model:

- Z-Normalized features, $\lambda = 10^{-5}$, no PCA

3.4 SVM Classifier

3.4.1 Linear SVM

Support Vector Machines are linear classifiers that look for maximum margin separation hyperplanes. The primal formulation of the soft-margin SVM problem consists in minimizing the function:

$$J(w, b) = \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N \max(0, 1 - z_i(w^T x_i + b))$$

where N is the number of training samples and C is an hyperparameter.

We are also going to take into account the dual formulation to solve the problem that consists in maximizing the function:

$$J^D(\alpha) = -\frac{1}{2} \alpha^T H \alpha + \alpha^T \mathbf{1} \text{ s.t. } 0 \leq \alpha_i \leq C \text{ and } \sum_{i=1}^N \alpha_i z_i = 0 \quad \forall i \in \{1, \dots, n\}$$

where $\mathbf{1}$ is a n -dimensional vector of ones and H is the matrix whose elements are $H_{ij} = z_i z_j x_i^T x_j$.

From the constraints of the Lagrangian problem that allows to introduce the dual formulation we can obtain that:

$$w^* = \sum_{i=1}^n \alpha_i^* z_i x_i$$

and the optimal bias b can be computed considering a sample x_i that lies on the margin: $z_i(w^{*T} x_i + b^*) = 1$. To be able to computationally solve the problem we need to modify the primal formulation as:

$$\hat{J}(\hat{w}) = \frac{1}{2} \|\hat{w}\|^2 + C \sum_{i=1}^N \max(0, 1 - z_i(\hat{w}^T \hat{x}_i))$$

where $\hat{x}_i = \begin{bmatrix} x_i \\ 1 \end{bmatrix}$ and $\hat{w} = \begin{bmatrix} w \\ b \end{bmatrix}$.

The scoring rule $\hat{w}^T \hat{x}_i = w^T x + b$ has the same form of the original formulation but we are also regularizing the norm of \hat{w} : $\|\hat{w}\|^2 = \|w\|^2 + b^2$ and we use a mapping $\hat{x}_i = \begin{bmatrix} x_i \\ K \end{bmatrix}$ to mitigate the fact that by regularizing the bias term we could obtain sub-optimal results. According to the modification done to the primal formulation we also modify the dual formulation as:

$$J^D(\alpha) = -\frac{1}{2}\alpha^T \hat{H}\alpha + \alpha^T \mathbf{1} \text{ s.t. } 0 \leq \alpha_i \leq C, \forall i \in \{1, \dots, n\}$$

where the equality constraint dissappeared (the one that L-BFGS was not able to incorporate) and the matrix \hat{H} can be computed as $\hat{H}_{i,j} = z_i z_j \hat{x}_i^T \hat{x}_j$. The




Figure 9: Linear SVM ($K=0$, Z-Normalized features) - minDCF for different values of C and different priors

Table 6: Linear SVM - 3-fold cross validation

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Z-normalized features - no PCA			
Linear SVM ($C = 0.1$)	0.158	0.052	0.141
Linear SVM ($C = 1$)	0.129	0.047	0.130
Z-normalized features - PCA(m=11)			
Linear SVM ($C = 0.1$)	0.275	0.102	0.233
Linear SVM ($C = 1$)	0.269	0.100	0.226
Z-normalized features - PCA(m=10)			
Linear SVM ($C = 0.1$)	0.295	0.114	0.259
Linear SVM ($C = 1$)	0.301	0.113	0.267

Table 7: Best models analyzed up to now

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Gaussian Models			
Tied Cov (Z-Norm, no PCA)	0.122	0.046	0.127
Tied Cov (Gau, PCA(m=10))	0.212	0.082	0.207
Logistic Regression Models			
LLR (Z-Norm, $\lambda = 10^{-6}$, no PCA)	0.132	0.047	0.126
QLR (Z-Norm, $\lambda = 10^{-5}$, no PCA)	0.153	0.052	0.142

only preprocessing step that has been considered is Z-normalization. As we can notice from figure 9 better

results are achieved for smaller values of the hyperparameter C so the choice of it is crucial.

- Best performances on the target application are reached for $C = 1$ combined with Z-normalized features and no PCA.
- PCA is not helpful in reducing the minDCF in any of the considered applications even if with PCA(m=11) there is no a high performance degradation for the target application.

Comparison: The results achieved by the model trained with Z-normalization, $C = 1$ and no PCA are closed to the the results obtained by the Linear Logistic Regression Model but slightly worse with respect to the results obtained by the Tied Gaussian model.

Selected Linear SVM Model:

Z-normalized features - K=0, C=1 - no PCA

3.4.2 Kernel SVM

SVMs allow for non-linear classification through an implicit expansion of the features in a higher dimensional space. In contrast with Quadratic Logistic Regression classifier we don't have to compute an explicit expansion of the features space, it is sufficient to be able to compute the scalar product between the expanded features: $k(x_1, x_2) = \phi(x_1)^T \phi(x_2)$ where k is the kernel function. We have to replace \hat{H} with $\hat{H} = z_i z_j k(z_1, z_2)$. We are going to implement two types of kernel: polynomial and rbf.

- Polynomial kernel of degree d :

$$k(x_1, x_2) = (x_1^T x_2 + c)^d$$

- Radial Basis Function kernel:

$$k(x_1, x_2) = e^{-\gamma \|x_1 - x_2\|^2}$$




Figure 10: Polynomial SVM ($K=1$, $c=1$, $d=2$, raw features) - minDCF for different values of C

The Polynomial Kernel SVM model was trained with raw features since the Z-Normalization step led to very poor results for almost all the values of the hyperparameter C . By looking at figure 10 we can realize that the choice of C is again crucial and best performances are obtained for $C \leq 10^{-2}$. For this reason i have chosen $C = 10^{-4}$ to show more detailed results.

Table 8: Polynomial Kernel SVM - $C = 10^{-4}$ - 3-fold cross validation

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Raw features - no PCA			
Poly SVM ($c=0, d=2$)	0.465	0.158	0.396
Poly SVM ($c=1, d=2$)	0.187	0.060	0.164
Poly SVM ($c=1, d=3$)	0.408	0.157	0.563
Raw features - PCA(m=11)			
Poly SVM ($c=0, d=2$)	0.553	0.199	0.487
Poly SVM ($c=1, d=2$)	0.210	0.062	0.158
Poly SVM ($c=1, d=3$)	0.452	0.183	0.537
Raw features - PCA(m=10)			
Poly SVM ($c=0, d=2$)	0.746	0.258	0.576
Poly SVM ($c=1, d=2$)	0.465	0.158	0.396
Poly SVM ($c=1, d=3$)	0.784	0.280	0.815

- The Polynomial SVM is able to achieve better results with respect to some versions of the Linear SVM. Best performances for the target application are achieved by the model trained with $c = 1$ and $d = 2$.
- The use of PCA(m=11) helps in reducing the $minDCF$ for the $\tilde{\pi} = 0.9$ application for the model trained with $c = 1$ and $d = 2$ and doesn't affect that much the target application.
- PCA(m=10) led to worse performances in all the considered applications.

The other kernel that was employed is the RBF kernel and these are the results achieved: In this case it was




Figure 11: RBF SVM ($K = 1, \gamma \in \{0.1, 0.01\}$) - $minDCF$ for different values of C

again used Z-normalization as pre-processing step. The hyperparameter γ has to be tuned so we trained the model for different values of γ to analyze the performances and here are showed the results for $\gamma = 0.1$ and $\gamma = 0.01$. Also the hyperparameter C is still to be choosed and needs to be estimated via cross-validation. The parameter K has been set to 1. From figure 11 we again realize that the value of the hyperparameter C is critical and we should choose a value $C \geq 10^{-1}$. The table below shows the results obtained with $C = 10$ and also results obtained with $\gamma = 1$ appear for completeness.

Table 9: RBF Kernel SVM - $C=10$

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Z-Normalized features - no PCA			
RBF SVM ($\gamma = 1$)	0.291	0.095	0.281
RBF SVM ($\gamma = 0.1$)	0.163	0.056	0.137
RBF SVM ($\gamma = 0.01$)	0.153	0.049	0.133
Z-Normalized features - PCA(m=11)			
RBF SVM ($\gamma = 1$)	0.362	0.123	0.350
RBF SVM ($\gamma = 0.1$)	0.255	0.092	0.255
RBF SVM ($\gamma = 0.01$)	0.254	0.091	0.211
Z-Normalized features - PCA(m=10)			
RBF SVM ($\gamma = 1$)	0.386	0.136	0.379
RBF SVM ($\gamma = 0.1$)	0.271	0.106	0.273
RBF SVM ($\gamma = 0.01$)	0.291	0.107	0.239

Table 10: Best models analyzed up to now

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Gaussian Models			
Tied Cov (Z-Norm, no PCA)	0.122	0.046	0.127
Tied Cov (Gau, PCA(m=10))	0.212	0.082	0.207
Logistic Regression Models			
LLR (Z-Norm, $\lambda = 10^{-6}$, no PCA)	0.132	0.047	0.126
QLR (Z-Norm, $\lambda = 10^{-5}$, no PCA)	0.153	0.052	0.142
SVM Models			
LSVM (Z-Norm, $C = 1$, no PCA)	0.129	0.047	0.130

- RBF kernel version of the SVM achieves results that are near to the ones obtained with the linear version and behaves better w.r.t. the Polynomial version
- The most promising value of γ seems to be $\gamma = 0.01$ because with this value we are able to obtain the best $minDCF$ for the target application (when no PCA is applied)
- PCA is not able to provide any better results in terms of $minDCF$

Comparison: The linear version of the SVM achieves better results with respect to the two kernel versions. This is aligned with the fact that also in Logistic Regression linear model worked better. **Selected Kernel SVM Models:**

- Polynomial SVM: Raw features, $K = 1, C = 10^{-4}, c = 1, d = 2, \text{PCA}(m=11)$
- RBF SVM: Z-Normalized features, $K = 1, C = 10, \gamma = 0.01, \text{no PCA}$

3.5 GMM Classifier

The last model we are going to take into account is a generative model. The GMM problem is related to the estimation of a population distribution and can

be also applied to the classification task. We resort the class-posterior probability used for the Gaussian classifiers:

$$f_X(x) = \sum_{c=1}^K f_{X|C}(x|c) \cdot P(C = c) = \sum_{c=1}^K \mathcal{N}(x|\mu_c, \Sigma_c) \cdot \pi_c$$

more in general we can introduce a weight term instead if the prior probability that will be one of the model paramers of the GMM:

$$X \sim GMM(M, \mathcal{S}, w) \Rightarrow f_X(x) = \sum_{g=1}^M w_g \cdot \mathcal{N}(x|\mu_g, \Sigma_g)$$

so the GMM density is the sum of M Gaussians where $M = [\mu_1, \dots, \mu_M]$, $\mathcal{S} = [\Sigma_1, \dots, \Sigma_M]$, $w = [w_1, \dots, w_M]$ are the model parameters. Gaussian components can be seen as clusters the samples belong to (in a hard or in a soft way) and the cluster label is a latent random variable. If we define

$$f_{X_i, G_i}(x_i, g) = w_g \mathcal{N}(x_i|\mu_g, \Sigma_g) \text{ and } f_{X_i}(x_i) = \sum_{g=1}^M f_{X_i, G_i}(x_i, g)$$

we can introduce a term called *responsability* that represents the posterior probability that a sample belongs to a certain cluster(component):

$$\gamma_{g,i} = P(G_i = g|X_i = x) = \frac{f_{X_i, G_i}(x_i, g)}{f_{X_i}(x_i)} = \frac{w_g \mathcal{N}(x_i|\mu_g, \Sigma_g)}{\sum_{g'} w_{g'} \mathcal{N}(x_i|\mu_{g'}, \Sigma_{g'})}$$

We can assign the sample to the cluster label for which the responsibility is maximum and then re-estimate the model parameters given the cluster assignments. An evident problem is that we are forming hard clusters so we are not admitting that a sample could belong to more than one cluster(component). To handle soft-clusters we introduce the statistics:

$$F_g = \sum_{i=1}^N \gamma_{g,i} x_i, \quad S_g = \sum_i \gamma_{g,i} x_i x_i^T, \quad N_g = \sum_{i=1}^N \gamma_{g,i}$$

to obtain a re-estimate of $w_g = \frac{N_g}{N}$. In this way we will be able to apply the Expectation-Maximization algorithm:

- E-step: estimate (given the model parameters $(M_t, \mathcal{S}_t, w_t)$):

$$\gamma_{g,i} = P(G_i = g|C_i = x, M_t, \mathcal{S}_t, w_t)$$

- M-step: estimate the new model parameters by using the statistics mentioned above

The estimation goes on starting from an initial value of the model parameters until a certain criterion is met. The EM algorithm thus require an initial estimate for the GMM parameters so we employ the LBG algorithm to incrementally construct a GMM with $2G$ components from a GMM with G components. The starting point will be $(1, \mu, C)$ so we use the empirical mean and covariance matrix of the dataset. We can

then build a 2-components model starting from a single one and from each of the new components we generate other new 2 components and so on and so forth.

We are also going to introduce a Diagonal and Tied version of the GMM problem: the diagonal variant consists of a model where each component has a diagonal covariance matrix (this does not correspond with the Naive Bayes assumption did for the Gaussian models). The Tied covariance model assumes that the covariance of a single GMM $\Sigma_{c,g} = \Sigma_c$ are the same but each GMM of each class has a difference covariance matrix Σ_c (again, this is different from the Tied MVG model).

We are now going to train a GMM over the samples of each class. Given the fact that almost all the features are already well distributed according to the Gaussian hypothesis i suppose that we won't need many gaussian components for the model. Moreover, Tied GMM model is supposed to outperform the other two models as already seen in the Gaussian classifiers. By considering




Figure 12: GMM - minDCF for different number of components

only models trained with 8 components we are now going to show more in detail the results obtained when training with and without PCA.

- As showed in figure 12, for a relative small number of components (8) good results are achieved by all the three types of models.
- Even if diag assumption doesn't hold really well especially when the number of components is low it performs better with higher number of components.
- The Tied model is the one with best results even if the Full model still performs really well and this corresponds to the hypothesis made before according to what has already been analyzed for the Gaussian models.
- Gaussianization never helps in achieving better results. According to what we had seen with the MVG models the already good distribution of data makes this classification task not benefitting from the Gaussianization pre-processing step and worse results are achieved when it is applied.
- PCA is not helpful in achieving better results (we can see an improvements for some applications when PCA(m=10) is applied w.r.t to PCA(m=11) but all these results are worse w.r.t the no PCA

Table 11: GMM - 3-fold cross validation

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Z-normalized features - no PCA			
GMM Full (8 comp.)	0.109	0.040	0.102
GMM Tied (8 comp.)	0.099	0.031	0.075
GMM Diag (8 comp.)	0.214	0.085	0.215
Z-normalized features - PCA(m=11)			
GMM Full (8 comp.)	0.229	0.074	0.175
GMM Tied (8 comp.)	0.198	0.067	0.166
GMM Diag (8 comp.)	0.311	0.111	0.275
Z-normalized features - PCA(m=10)			
GMM Full (8 comp.)	0.225	0.082	0.212
GMM Tied (8 comp.)	0.204	0.072	0.186
GMM Diag (8 comp.)	0.310	0.108	0.260
Gaussianized features - no PCA			
GMM Full (8 comp.)	0.221	0.076	0.210
GMM Tied (8 comp.)	0.158	0.059	0.151
GMM Diag (8 comp.)	0.419	0.162	0.397
Gaussianized features - PCA(m=11)			
GMM Full (8 comp.)	0.246	0.091	0.199
GMM Tied (8 comp.)	0.175	0.066	0.169
GMM Diag (8 comp.)	0.477	0.174	0.415
Gaussianized features - PCA(m=10)			
GMM Full (8 comp.)	0.274	0.096	0.257
GMM Tied (8 comp.)	0.185	0.069	0.185
GMM Diag (8 comp.)	0.406	0.148	0.356

versions of the trained models). With Gaussianization and PCA(m=11) instead we can notice an improvement for $\tilde{\pi} = 0.9$ application.

Table 12: Best models analyzed up to now

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Gaussian Models			
Tied Cov (Z-Norm, no PCA)	0.122	0.046	0.127
Tied Cov (Gau, PCA(m=10))	0.212	0.082	0.207
Logistic Regression Models			
LLR (Z-Norm, $\lambda = 10^{-6}$, no PCA)	0.132	0.047	0.126
QLR (Z-Norm, $\lambda = 10^{-5}$, no PCA)	0.153	0.052	0.142
SVM Models			
LSVM (Z-Norm, no PCA)	0.129	0.047	0.130
Poly SVM (Raw, $c = 1$, $d = 2$, PCA(m=11))	0.210	0.062	0.158
RBF SVM (Z-Norm, $\gamma = 0.01$, no PCA)	0.153	0.049	0.133

The GMM Tied(8 components) trained with Z-Normalization and no PCA is the one that achieves the best results among all the models evaluated up to now.

Selected GMM Model:

- GMM Tied (8 componenets), Z-Normalization, no

PCA

4 Score Calibration

4.1 Calibration Analysis On Selected Models

We now select one candidate for each of the previous analyzed classification models (Gaussian, Logisti Regression, SVM, GMM):

- Gaussian model: Tied-Cov (Z-Normalization, no PCA)
- Logistic Regression: Linear LR (Z-Normalization, no PCA)
- SVM: Linear SVM (Z-Normalization, no PCA)
- GMM: Tied GMM with 8 components (Z-normalization, no PCA)

4.2 Calbrating Scores For Selected Models

We are going to transform the scores provided by each model so that the theoretical threshold $t = -\log \frac{\tilde{\pi}}{1-\tilde{\pi}}$ provides close to optimal values over a wide range of effective priors $\tilde{\pi}$. What we want to find is a monotonic function f that maps not-calibrated scores in calibrated scores.

We assume that the function f has the form:

$$f(s) = \alpha s + \beta$$

and $f(s)$ can be interpreted as log-likelihood ratio for the two classes hypotheses:

$$f(s) = \log \frac{f_{S|C}(s|\mathcal{H}_T)}{f_{S|C}(s|\mathcal{H}_F)} = \alpha s + \beta$$

and the class posterior probability for prior $\tilde{\pi}$ corresponds to:

$$\log \frac{P(C=\mathcal{H}_T|s)}{P(C=\mathcal{H}_F|s)} = \alpha s + \beta + \log \frac{\tilde{\pi}}{1-\tilde{\pi}}$$

By interpreting scores s as samples of a dataset (where each sample has 1 feature) we can employ a prior weighted logistic regression model to learn the model parameters over our calibration set (the dataset composed of the scores for a certain model). If we let:

$$\beta' = \beta + \log \frac{\tilde{\pi}}{1-\tilde{\pi}}$$

we have exactly a Logistic Regression model where α and β' are the model parameters we have to learn. We have still to specify a prior $\tilde{\pi}$ that will be the one of our target application even if we will notice that the improvements in terms of calibration will also involve the unbalanced applications. To obtain calibrated scores we will have to compute:

$$f(s) = \alpha s + \beta = \alpha s + \beta' - \log \frac{\tilde{\pi}}{1-\tilde{\pi}}$$




Figure 13: Bayesian Error Plots for each of the selected models (without score calibration)

The Tied Gaussian, Linear Logistic Regression and Tied GMM models seem to be already well calibrated in particular for positive values of the prior log-odds. The Tied Gaussian model could be slightly improved in terms of calibration for negative values of the prior log odds. The SVM instead is not well calibrated for almost every considered values of the prior log-odds.




Figure 14: Bayesian Error Plots for each of the selected models (with score calibration)

Table 13: Calibrated vs not calibrated scores for best selected models

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Uncalibrated scores [minDCF / DCF]			
Tied Gaussian	0.122/0.127	0.046/0.050	0.127/0.134
Lnear LR	0.132/0.139	0.047/0.048	0.126/0.136
Linear SVM	0.132/0.289	0.047/0.047	0.129/0.285
Tied GMM	0.099/0.110	0.031/0.031	0.075/0.081
Calibrated scores [minDCF / DCF]			
Tied Gaussian	0.122/ 0.128	0.046/ 0.049	0.127/ 0.135
Lnear LR	0.132/ 0.141	0.047/0.048	0.126/ 0.132
Linear SVM	0.132/ 0.135	0.047/0.047	0.129/ 0.137
Tied GMM	0.099/ 0.105	0.031/0.031	0.075/ 0.076

We can easily notice that for models where scores were already well calibrated the score calibration didn't provide effective benefits (in some cases we also obtained slightly worse results). For SVM model great benefits are obtained for the unbalanced applications.

5 Experimental Results

Now we are going to train all the previous models on the entire training set and the evaluation set will be used for the first time to assess if the results obtained so far are consistent and the best model selected is

confirmed (we could also realize that other choices would have been better for this evaluation set).

Table 14: minDCF over evaluation set for all models trained over the whole training set (Z-Normalized features, no PCA)

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Gaussian Models (Z-normalized features, no PCA)			
Full Cov	0.134	0.053	0.138
Tied Cov	0.133	0.051	0.135
Naive Bayes	0.810	0.570	0.882
LR Models (Z-Normalized features, no PCA)			
LLR ($\lambda = 10^{-6}$)	0.135	0.052	0.133
QLR ($\lambda = 10^{-5}$)	0.148	0.053	0.141
SVM Models (Z-Normalized features, no PCA)			
LSVM ($K = 0, C = 1$)	0.142	0.052	0.137
Poly SVM ($K = 1, C = 10^{-4}, c = 1, d = 2$)	0.194	0.062	0.163
RBF SVM ($K = 1, C = 10, \gamma = 0.01$)	0.138	0.054	0.133
GMM Models (Z-Normalized features, no PCA)			
GMM Full (8 comp.)	0.092	0.034	0.086
GMM Full (16 comp.)	0.103	0.043	0.117
GMM Tied (8 comp.)	0.088	0.030	0.087
GMM Tied (16 comp.)	0.089	0.033	0.086
GMM Diag (8 comp.)	0.236	0.090	0.234
GMM Diag (16 comp.)	0.244	0.098	0.232

Table 15: minDCF over evaluation set for all models trained over the whole training set (Z-Normalized features, PCA($m=11$))

	$\tilde{\pi} = 0.1$	$\tilde{\pi} = 0.5$	$\tilde{\pi} = 0.9$
Gaussian Models (Z-normalized features, PCA($m=11$))			
Full Cov	0.280	0.113	0.255
Tied Cov	0.278	0.111	0.256
Naive Bayes	0.293	0.116	0.263
LR Models (Z-Normalized features, PCA($m=11$))			
LLR ($\lambda = 10^{-6}$)	0.281	0.110	0.259
QLR ($\lambda = 10^{-5}$)	0.253	0.108	0.258
SVM Models (Z-Normalized features, PCA($m=11$))			
LSVM ($K = 0, C = 1$)	0.277	0.110	0.262
Poly SVM ($K = 1, C = 10^{-4}, c = 1, d = 2$)	0.189	0.068	0.151
RBF SVM ($K = 1, C = 10, \gamma = 0.01$)	0.243	0.109	0.246
GMM Models (Z-Normalized features, PCA($m=11$))			
GMM Full (8 comp.)	0.179	0.080	0.196
GMM Full (16 comp.)	0.219	0.086	0.211
GMM Tied (8 comp.)	0.173	0.072	0.195
GMM Tied (16 comp.)	0.185	0.078	0.194
GMM Diag (8 comp.)	0.329	0.124	0.312
GMM Diag (16 comp.)	0.311	0.141	0.327

The results obtained by training the models on the

entire training set and by evaluating them on the entire evaluation set are coherent with the results previously obtained and with the observations made. The ROC curve in figure 15 shows (and confirms) that the best performing model is the Tied GMM with 8 components. The other three models also perform really well but slightly worse than this.




Figure 15: ROC curves of selected models (on the right a zoomed version)

6 Conclusions

We can conclude by saying that in general linear models have better performances on these data with respect to the non-linear ones. Due to the high level of correlation among features the PCA didn't help in achieving better results (despite the fact that with PCA($m=11$) we have obtained not a big performance degradation in most of the models) and Naive hypothesis didn't worked well in Gaussian models. We have been able to reach a $\min DCF \approx 0.03$ for the target application($\tilde{\pi} = 0.5$) but also for the unbalanced applications we were able to achieve good results ($\min DCF \approx 0.1$ for $\tilde{\pi} = 0.1$ and $\min DCF \approx 0.1$ for $\tilde{\pi} = 0.9$). The choices made on the training set via k-fold cross-validation proved to be coherent with the results obtained on the entire training set.